

# Shorter ideas. Complete books. Room for the artwork.

## Codex implementation patch and authoring-policy update

Mario Sinclair | October 4, 2026 (America/New\_York)

This supersedes conflicting image-count, idea-count, and lesson-duration guidance in Handoff 002 and the prepared companion bundles. Preserve the accepted reading experience and manual Files picker. This is a specification for the Mac Codex engineer, not an applied source-code patch or an importable book.

## The decisions

- **1-12 ideas per new book collection.** Twelve is a ceiling, never a quota. Select fewer when that teaches the material well. Deliver every selected idea together, including on updates.
- **No numerical image-count ceiling.** Remove the 32-asset rule without replacing it with another count. Keep byte, dimension, type, and decoding protections.
- **Five minutes is the upper editorial limit for a complete idea.** Include introduction, main idea/story, explanation, takeaway, transitions, both questions and choices, feedback, and a realistic answer allowance. Shorter is welcome; do not pad to four minutes.
- **Keep manual import now.** GitHub Releases plus a small catalog is the agreed future distribution direction. No cloud browsing, accounts, downloading service, or publishing is part of this patch.

## Baseline to inspect

Repository: [Patchagray/LifeIsLearned](#)

Reviewed branch: **feature/handoff-002-ui-polish**

Reviewed commit: d10a60f84b9921cd24a74666b5877359f0623090

The default branch previously pointed to an older bootstrap version. Inspect the current remote and preserve any newer work. Branch from the accepted implementation; do not build this patch on the old bootstrap branch. Follow AGENTS.md; do not change signing or merge into the default branch.

# Collection and artwork rules

## Apply the cap at the correct boundary

Fresh imports and updates must contain 1-12 unique ideas, with an ordered manifest matching all lesson IDs and revisions. Reject 13 or more before preview/publication, with an actionable message: "This collection contains 13 ideas. Prepare a complete release with no more than 12 selected ideas." Do not auto-truncate or split a book into disguised volumes to bypass the cap.

Keep already-installed, previously valid collections readable, including format-2 collections with 13-100 ideas. Do not let this policy change trigger recovery mode, hide books, or discard progress on launch. Separate stored-content validation from new-import policy; do not merely lower a shared limit everywhere. An update to an over-cap installed book must meet the new cap and explicitly declare removals for learner review. No silent migration to 12.

## Remove count, retain resource protection

Remove asset-count rejection from native import validation and the Python authoring validator. Keep the existing limits: **64 MiB encoded package, 2 MiB per image, 24 MiB total image-file bytes after base64 decoding, 2048 x 2048 maximum dimensions, one image frame, PNG/JPEG type and native-decode checks**. These are resource protections, not an image-count limit. Unique nonempty asset IDs and valid image references remain required.

Continue storing each shared image once in assets and referring to it by imageID; include the cover in the same resource budgets. Optimize illustrations rather than weakening byte limits. No dependency, network image fetch, new package container, or formatVersion bump is needed solely for these policy changes.

## Protect revisions and progress

Keep existing fullCollection, manifest, collectionRevision, stable idea IDs, explicit removedLessonIDs, import preview, atomic publication, no-op detection, and archived progress behavior. Changed narration, quizzes, estimates, referenced sources, or artwork require affected idea revision increases under the existing comparison rules. Changed releases require a collection revision increase.

When ideas are merged, do not transfer completed status onto unpracticed material. Retain an ID only if it still represents the same core idea; otherwise create a new stable ID and explicitly retire the replaced IDs. Keep all old progress archived.

# A five-minute whole-idea budget

## Count what the learner actually hears

Create a repeatable authoring timing report. Count narrated page titles and prose, question prompts, option-number labels and all choices, one feedback response per question, spoken feedback prefixes, and the completion announcement. Count a book preface too if the implementation speaks it in the first-idea journey. Do not count sources or interface text unless they are actually narrated.

**Planning formula:** spoken-word time + configured transition/reflection pauses + 40 seconds of answering time (20 seconds for each of two questions). For conservative budgeting, use the longest single feedback choice for each question. Start at 130 words/minute, then verify with the actual reference voices at normal speed and the fresh-install 2-second transition setting. Record those settings and actual durations.

## Draft shorter than the ceiling

Start around **350-450 total spoken words**, with an original example/story around **150-220 words**. These are drafting guides, not new quotas or guarantees. A 450-word script takes about 208 seconds at 130 wpm; adding eight 2-second transitions and 40 seconds for answers gives about 264 seconds. Extra titles, pauses, feedback, or slower voices consume the remaining margin. Shorten further when measurement requires it.

Keep one teachable distinction, a concise original example, a short explanation/takeaway, and exactly two meaningful application questions for the authored releases. Use fewer story screens where useful; retain the valid intro-to-takeaway structure. Do not make prose unnaturally dense or speed up narration to pass.

## Release gate, not a countdown

The reference first-pass journey must total **300 seconds or less**, including the answer allowance. A 4-5 minute story alone fails. Produce a per-idea timing report; block authoring approval/export when over budget. An estimatedMinutes value of 5 is not proof of compliance. Display an approximate whole-idea duration rather than a story-only estimate; calculate it consistently.

Never enforce the budget by cutting speech, skipping feedback, auto-answering, forcing navigation, or reducing a user-selected pause. Slower voices/speeds, longer deliberation, optional replays, and retries can extend actual elapsed time. Disclose that distinction in the duration help text. Preserve accessibility and user control. Runtime import validation can check structure, but cannot prove a learner will finish within five minutes.

# Editorial selection and existing books

## Explain the selection honestly

Use a short overview preface, visible before starting. Suggested copy: **"This companion presents [N] main ideas we selected from the publicly available material we reviewed. The book contains additional ideas and detail; this is a curated selection, not an exhaustive summary or the author's own ranking."**

Replace [N] with the actual count; never show a placeholder to readers.

State the actual coverage and sources in coverageNote. If the full book was not reviewed, say so. Select for centrality in reviewed material, usefulness, distinctness, and teachability; retain meaningful source order where supported. Do not call these "the author's top 12" without an explicit source establishing that ranking. Publicly accessible source material does not itself grant redistribution rights; continue using original text/art with attribution.

## Both prepared collections need editing

- **The Influential Mind:** ten ideas already satisfy the count cap, but all must be shortened and retimed as complete experiences. Its prior four-minute stories did not include teaching and quizzes. Keep the established book ID influential-mind. Reconcile priors with the installed revision-1 demo and increment changed content appropriately.
- **Never Split the Difference:** reduce the earlier fourteen-idea draft to at most twelve through a documented editorial selection or sensible consolidation, then shorten every idea. Do not merely delete the final two ideas or compress several unrelated concepts into a rushed lesson.
- **Starter/demo:** shorten the accepted test idea using the same timing budget, preserving voices, meaningful feedback, and source honesty. Increase its revision for the changed story/questions. Make the revised demo available as a reviewed full-collection update; do not silently replace an installed user collection through a startup seed change.

Treat the existing companion ZIPs as superseded drafts for final packaging. Do not label them ready to import. Their artwork briefs may be reused only after checking fit with the shortened stories. No image generation or bulk book import is required to complete the engineering patch.

## Updated reusable authoring instruction

"Prepare one complete formatVersion-2 collection with 1-12 selected ideas and an honest coverage preface. Teach one clear idea per lesson with concise original examples, guide/storyteller roles, and exactly two application questions with individual feedback. Keep the entire reference first-pass idea at or below five minutes, including all narration, transitions, quiz choices, feedback and an answer allowance. Shorter is welcome. Include a timing report, stable IDs/revisions, shared original artwork, source references and the complete ordered manifest. Never invent unavailable book content or imply the author ranked our selection."

# Implementation map and acceptance

## Files and responsibilities

**Models/LessonPackage.swift:** separate new-import policy from stored-content readability; enforce 12 for new imports and remove asset-count rule.

**Models/CollectionContract.swift; Services/CollectionStorage.swift:** keep resource validation, staging, revision comparisons and recovery coherent.

**Tools/validate\_package.py:** match new import policy; add repeatable whole-idea timing reporting/gating for authored releases.

**Tools/convert\_package.py and fixture generators:** do not export over-cap releases or use 100 ideas as a normal valid new-import fixture.

**Tests/CollectionTests.swift and related tests:** update boundary/stress fixtures; preserve real over-cap stored-library migration cases.

**CONTENT\_AUTHORING.md, AGENTS.md, VALIDATION.md, PERSISTENCE.md and relevant UI**

**duration copy:** synchronize policy and document legacy handling. Models/ and Services/ paths are relative to LifeIsLearned/. Tests/, Tools/ and document paths are relative to the repository root.

## Required evidence

- Accept new collections with 1 and 12 ideas; reject 0 and 13. Preserve manifest and duplicate-ID checks. Load a previously valid stored 100-idea collection without recovery mode or truncation.
- Accept and relaunch a 12-idea fixture with at least 43 valid shared images within byte limits. Verify every reference and a cover. Reject an over-byte-budget package, oversize dimension, malformed image and missing reference. Run this on iPhone and iPad; record responsiveness and memory observations.
- Keep existing no-op, changed-without-revision, downgrade, explicit-removal, unchanged-progress, revised-progress and failed-write recovery tests. Verify an over-cap stored book can receive a compliant full update with acknowledged removals.
- Check the timing report includes options, feedback and pauses. Demonstrate  $\leq 300$  seconds passes and  $> 300$  fails the authoring gate. Audition the shortened demo end to end; record guide/storyteller selections, speed, pause, total speech time and quiz allowance. Do not claim that a word-count estimate is measured playback.
- Build the relevant target, run affected tests and verify narration, manual controls, reflection boundaries, quiz retries, saved position and manual file import. Preserve existing signing and user data.

## Delivery and deferred work

Commit implementation, concise migration notes, timing evidence and test results to a dedicated branch; push and identify the exact commit. Do not merge automatically. Separate checks actually run from checks still pending. No cloud deployment now. Later: GitHub release assets plus a small catalog, complete-book background downloads, progress/cancel/retry, validation before installation, offline reading, and storage removal that preserves learning progress and actually reclaims old book files. Keep credentials out of the app.